Why should I take this course? ①

Why should I care about cloud?

virtualization?

rust?

1

what is cloud? (1.5)

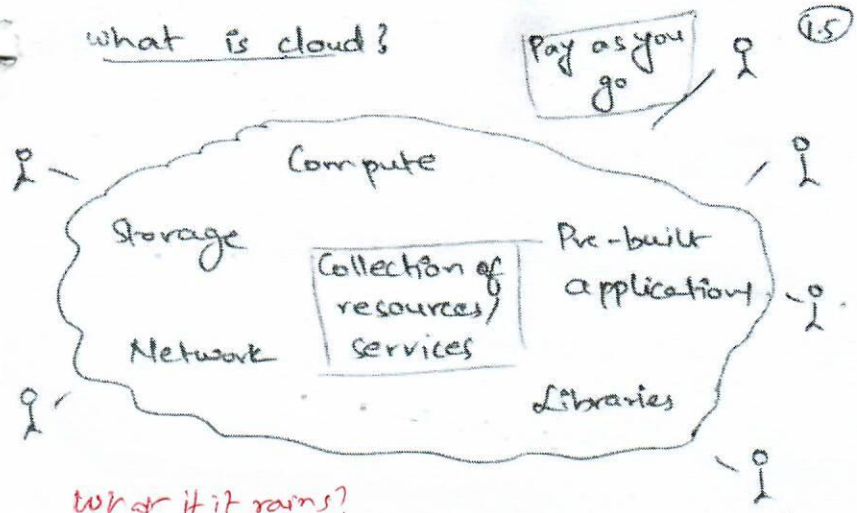

what if it rains?

why cloud? ①

Don't want to manage resources

☹ Hardware faults are common

Assume your laptop crashes after 5 years

≈ 1800 days

If you've 200 of them, one is crashing 9 days

Why cloud? Elasticity 2m

- My ~~blog~~ memes app today got very popular
  100 → 100,000 users
-  Laptop crashed. my chance of becoming popular is missed.
- Solution: use 1,000 laptops. · High upfront work

Why cloud? Provides fault tolerance 3m

- Doesn't lose data even when individual disks might fail
- Moves your computation to other machine if CPU overheats and dies.

→ AWS SLA ←

- Downtime is expensive > $1M/hour ①
  - Crashed server
  - Disconnected wifi/electricity
  - Tornado/Earthquake
- Loss of trust
  - Bank lost last 100 transactions

Why cloud? **Proximity** 1m

My memes app is getting popular in Canada.

- **Ping time ≈ 300 ms**
- Every 100 ms delay drops **7% users**
- Rent a VM next door to user ≈ 20ms

→ Region-wise ping times ←

Virtualization ①

What?
- Act of "virtualizing" physical resources

Why?
relavance to cloud

Challenges?

---

③

Investment 

20yrs sell

- Management, security overhead

Buy virtual gold: Gold bonds

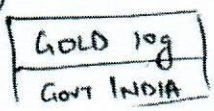

back to VMs ②

Why? (Buyer)

- No management
- Little upfront cost
- Elasticity
- Trusted

- DDos attac
- Power outage
- Earthquak
- SW update

Why? (Seller) ②

- Overprovision

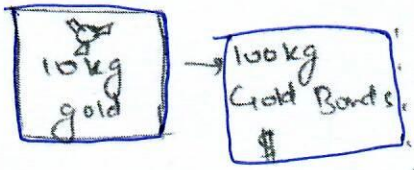

- Anyways buyi physical resources for self
- Amazon, Google
- Rent unused.
- Incremental mgmt overhe

Virtualization supports cloud (1) ②

- Live migration / checkpoint-restore
  - zero downtime

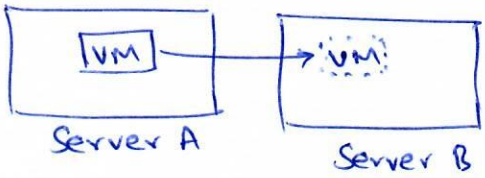

- Load balancing
- Management: restart/upgrade server A (supercomputers)
- Proximity.

* Strong Isolation (even performance)

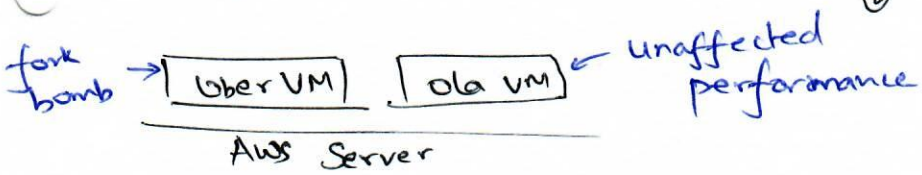

* Auditting / monitoring / Intrusion analysis
  - What all did the VM do?

State machine replication

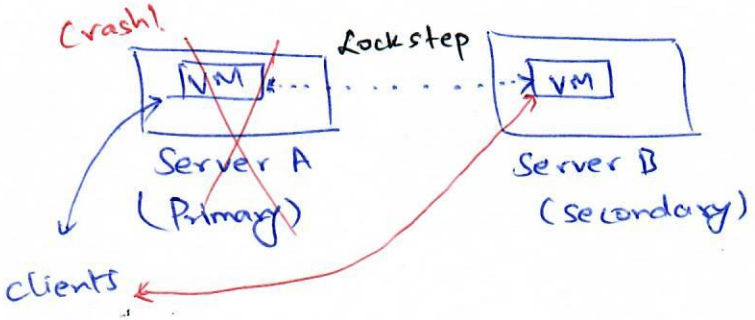

Fault tolerance!

- Binary ~~form~~ one architecture can run on another

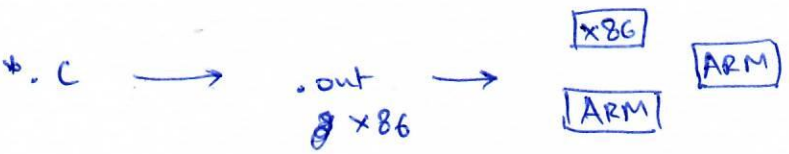

- Diversified placement
- Support legacy software

Virtualization challenges: Feasability ②

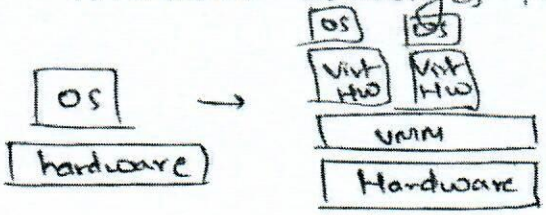

- Zero/minimal modification to OS, apps (OS thinks it owns HW)
- Performance

What's new? ⑫

Traditional IaaS - Rent VMs

1 How many VM?

Dropping requests ← Too few

Idle ⇒ losing money → Too many

Sophisticated capacity planning, elasticity

New cloud. ①

FaaS: specify functions, triggers

: Auto scaling

: Do not pay when function is not running >AWS Pricing<

New challenges: fast boot, dynamic provisioning, sharing data, etc

Pricing comparison (3)

* FaaS. Mobile application backend

  $2.33

* IaaS. a1.medium → 0.0255/hr * 24 * 30

  = $18.36

* Ignoring data transfer/storage costs

What's new (2) (2)

Monolithic → Microservices

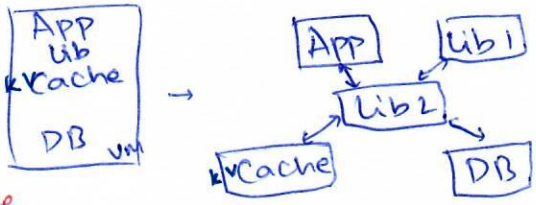

Can we improve performance? Specialize for one workload?

- Scale independently
- Different fault tolerance properties ⇒ crash Lib, can't crash DB

Why take this course? (1.5)

1. You'll work with cloud
   - Startup: using cloud infra (personal projects)
   - AWS / Google / Azure: Managing
   - ML experiments
   - Research: interesting problems
2. Hands on — virtualization: beautiful history
3. Study real systems → firecracker

Highly collaborative project (1)

- Collaboration is hard

=> GIT, CI/CD, PRs, unit tests

- Design is hard :) Interfaces/abstraction
  but fun

- Real end-to-end product

---

why rust? (2)

- Fastest growing language "most-loved" according to stack overflow
- Fast. Systems language C/C++
- Automatic memory management Java/Python/Go (without GC)
- Modern paradigms: Iterators, closures
- Safe: no null pointers, no double mutation, no pointer arithmetic