

✓ Last lecture: Memory virtualization ②

Transparency: I think I own memory

Base & bound:

addr translation → adder
+ check bounds

✓ Perf

✓ Hardware complexity

✗ Flexibility
✗ Utilization

3

✓ Paging ②

• Bits: valid, present, dirty, r/o

• Page table

• Large ⇒ Hierarchical

• TLB - *spatial locality* for TLB hits

• Demand paging

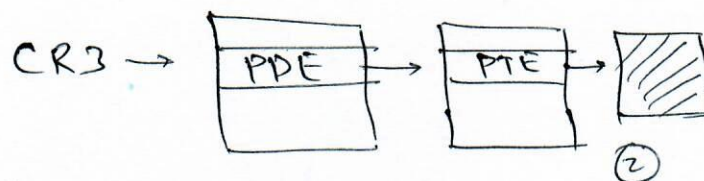
Perf ✓

HW complexity ✗

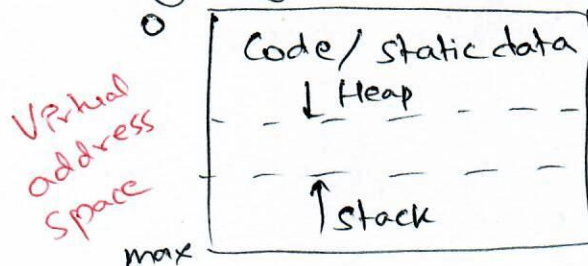
Flexibility ✓

Resource util ✓

✓ HPT



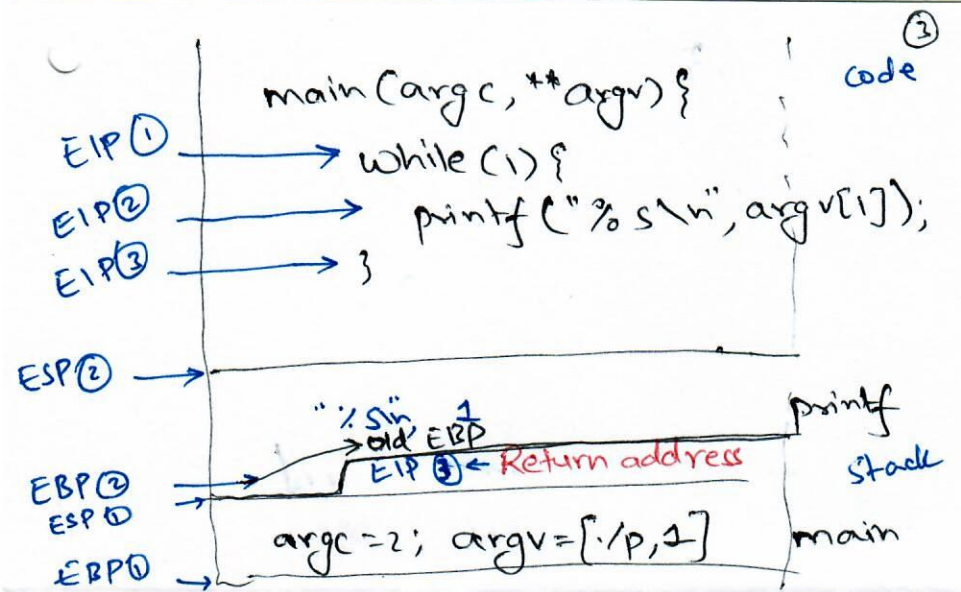
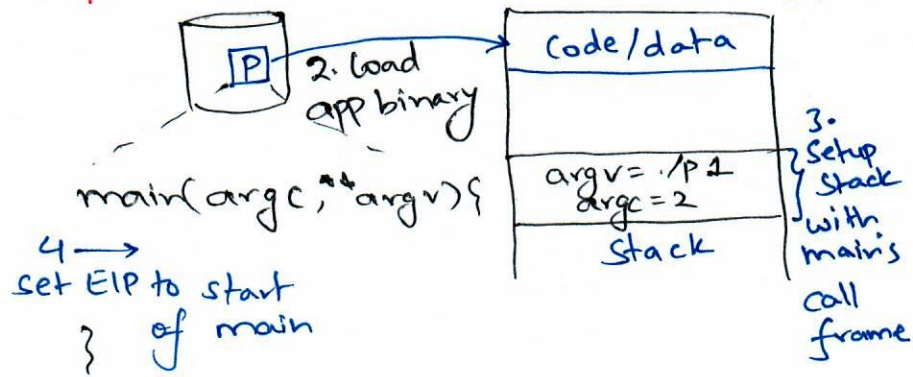
Memory layout



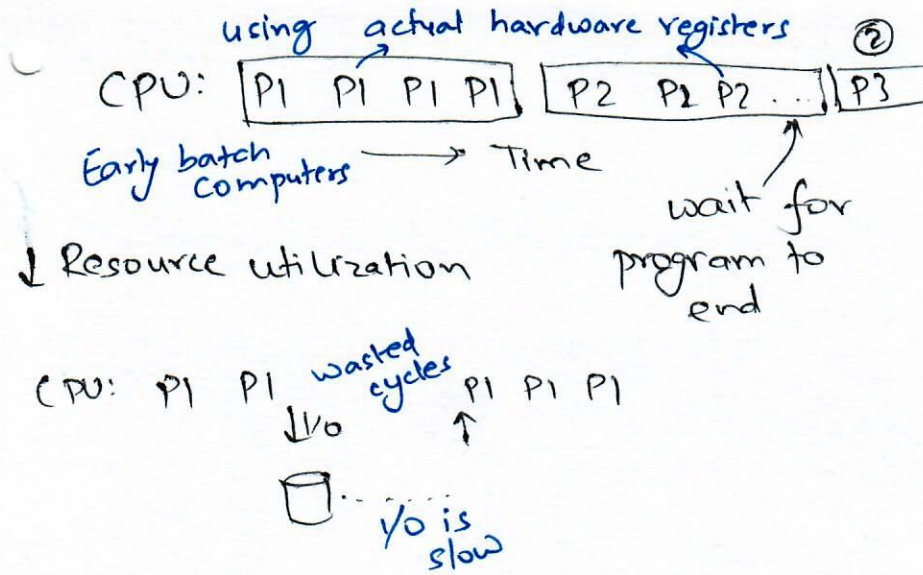
How does a program run? ③

\$./p1

1. New VA is allocated



- ②
- CPU register (x86)
- CR3 → Page table ← Hidden from programs
- EIP → Instruction pointer ← Auto Incremented by CPU
- managed by Compiler {
- ESP → Stack pointer. Top of stack
- EBP → Bottom of current call frame
- Where to allocate call frames
- General: eax, ebx, ecx...



- Operator (manual): ①
- Kill run away process
 - Prioritize process
 - manage list
 - Only used during office hours
 - Only way to intervene is to kill a process

P: `main(argc, **argv) {` ①

modifies actual hardware registers

`while (1)`

`printf("%.5f\n", argv[1]);`

`$./p A & ./p B & ./p C & ...`

A
B
A
C
B

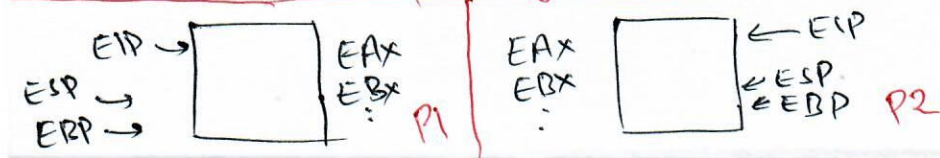
multiple processes are able to run even though there is single CPU

CPU virtualization:

②

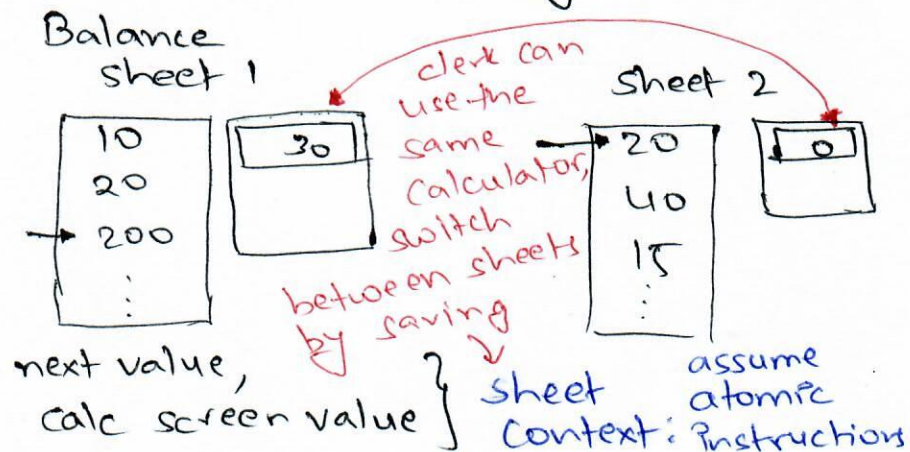
Every process THINKS that it owns CPU registers and directly changes them when it runs!

Transparency



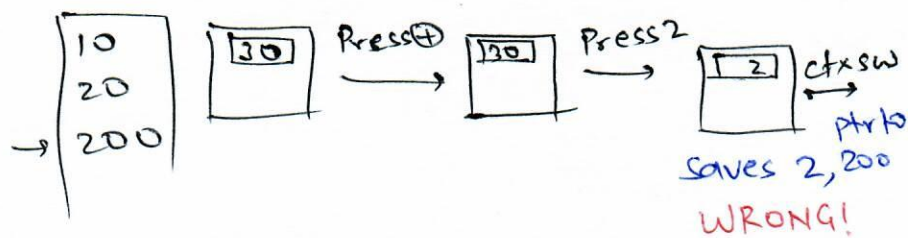
Bank clerk analogy

①



Context sw can't happen in middle of instruction

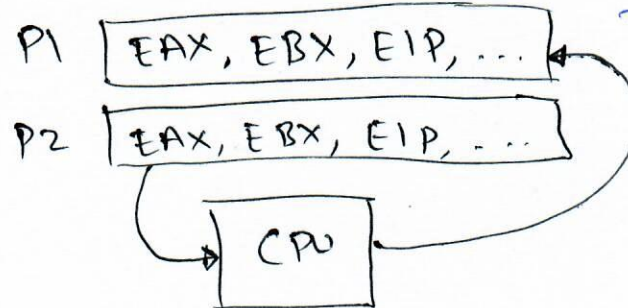
②



Sheet context: assumes 1 instruction → adding num
Each instruction → ATOMIC

Context switch

Process Control block



Can't happen in middle of instruction.

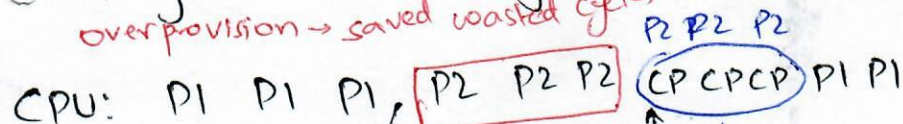
Comparison w/ operator

Flexibility ↑

- Run multiple processes simultaneously
- Insert high-priority job without rebooting
- Dynamically kill/spawn processes
- No need for operator → OS!

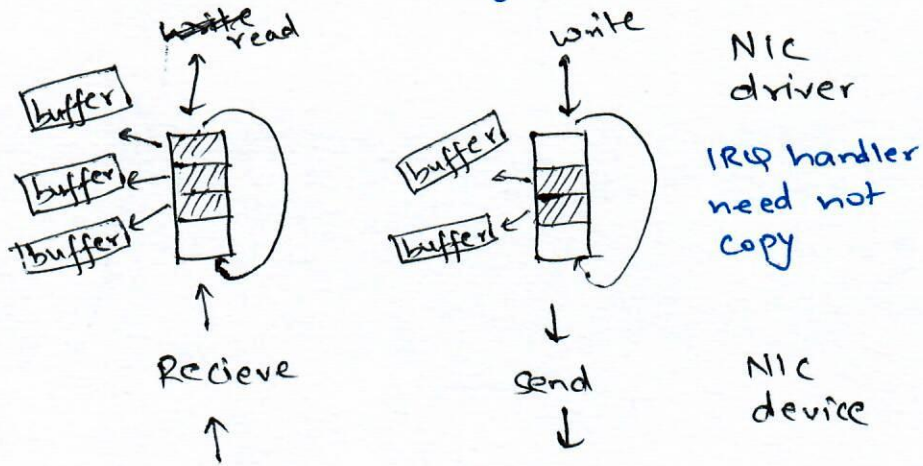
Performance benefit

overprovision → saved wasted cycles



Improved resource utilization

Aside: DMA ring buffers ②



Problem: Context switch is light ②

→ few registers

Impact on performance is bad.

$P1 \leftrightarrow P2$: different address spaces

D-cache, I-cache, TLB misses

↑ flushed in some architectures

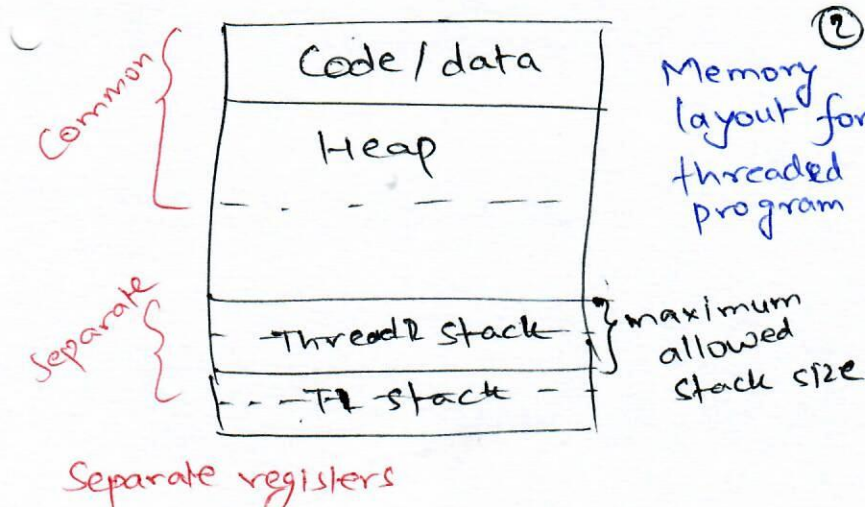
Improve perf. ①

- Only change registers
- Don't change address space
 - Threads

• T1 T1 T1 T2 T2 T2 T1 T1 T1

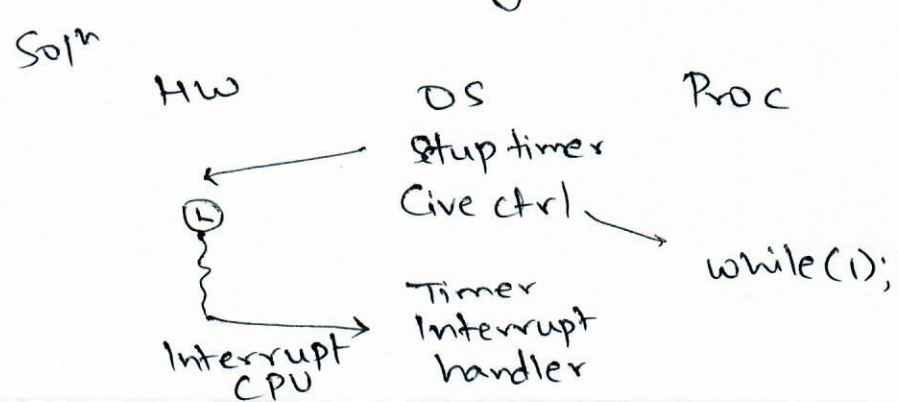
↑
fast

- very effective for I/O heavy processes



Prob Give control, keep ability to take back control

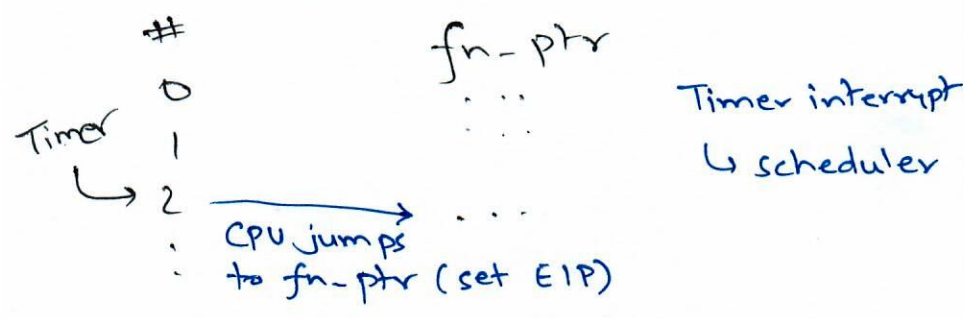
while (1); Single CPU



which handler?


Interrupt descriptor table

Set up at boot



Modify IDT

(2)

fn-ptr
0
1
→ 2  does not call sched

Privileged instruction ⇒ Traps to OS
⇒ Kill proc.

x86 rings

(1)

ring 3	User

_____	Other privileged inst ⁿ :
_____	- Disable interrupts
_____	- Change CR3
ring 0	kernel

Interesting programs need to use privileged instructions (2)

- Sending data on network
- Reading disk contents (100x heavier than function calls)

⇒ System calls

execute INT N instruction

↳ IDT